

Universidad Nacional del Sur

PROYECTO FINAL DE CARRERA
INGENIERÍA EN COMPUTACIÓN

*Seguridad en redes LAN: utilizando docker para
mejorar la infraestructura.*

Salvador Catalfamo

BAHÍA BLANCA – ARGENTINA
2020

Universidad Nacional del Sur

PROYECTO FINAL DE CARRERA
INGENIERÍA EN COMPUTACIÓN

*Seguridad en redes LAN: utilizando docker para
mejorar la infraestructura.*

Salvador Catalfamo

BAHÍA BLANCA – ARGENTINA
2020

Resumen

A lo largo de la carrera, hemos visto como las organizaciones abordan los temas de seguridad en sus sistemas informáticos. Mayormente, se concentran en los equipos que están expuestos a la red pública, dejando de lado los que se encuentran aislados de la misma. Erróneamente, muchas veces se piensa que es suficiente, sin embargo, puede traer graves inconvenientes. Es por eso que realizaremos un estudio teórico/práctico sobre las consecuencias de la navegación en redes internas sin ningún tipo de cifrado de datos ni certificaciones.

PALABRAS CLAVE:

Seguridad e Infraestructura

Docker

Linux

Kali

Máquinas Virtuales

Índice general

1. Introducción	1
1.1. Objetivos	1
1.2. Plan de tesis y principales contribuciones	2
1.3. Trabajos previos relacionados	2
2. (Ajustar) ¿Qué circula por una red interna?	1
2.1. Introducción	1
2.2. Protocolos asociados a la web	1
2.2.1. ¿Que es el protocolo HTTP?	1
2.2.2. Arquitectura	2
2.2.2.1. Client/Server Messaging	2
2.2.2.2. Ejemplo	3
2.2.3. Formato del mensaje (Mejorar intro)	3
2.2.3.1. Start Line	4
2.2.3.1.1. Request Line	4
2.2.3.1.2. Status Line	4
2.2.3.2. Header Fields	5
2.2.3.2.1. Field Extensibility	5
2.2.3.2.2. Field Order	5
2.2.3.2.3. Whitespace	5
2.2.3.2.4. Field Parsing	6
2.2.3.2.5. Field Limits	6
2.2.3.2.6. Field Value Components	6
2.2.3.3. Message Body	7
2.2.3.3.1. Transfer-Encoding	8

2.2.3.3.2.	Content-Length	8
2.2.3.3.3.	Message Body Length	9
2.2.4.	Métodos del protocolo HTTP	11
2.2.5.	Response Status Codes	14
2.2.6.	Overview of Status Codes	15
2.2.6.1.	Informational 1xx	17
2.2.6.2.	Successful 2xx	17
2.2.6.3.	Redirection 3xx	17
2.2.6.4.	Client Error 4xx	18
2.2.6.5.	Server Error 5xx	18
2.2.7.	HTTPS con SSL	18
2.2.7.1.	SSL Roles	18
2.2.7.2.	SSL Messages	19
2.2.7.2.1.	ClientHello	19
2.2.7.2.2.	ServerHello	20
2.2.7.2.3.	ServerKeyExchange	20
2.2.7.2.4.	ServerHelloDone	21
2.2.7.2.5.	ClientKeyExchange	21
2.2.7.2.6.	ChangeCipherSpec	21
2.2.7.2.7.	Finished	22
2.2.7.2.8.	Ending Secure Communications	23
2.2.7.3.	Authenticating the Server's Identity	23
2.2.7.3.1.	Certificate	24
2.2.7.3.2.	ClientKeyExchange	24
2.2.7.4.	Certificate Functionality	25
2.2.7.4.1.	Single Domain	25
2.2.7.4.2.	Multi-Domain	25
2.2.7.4.3.	Wildcard	25
2.2.7.4.4.	Multi-Domain Wildcard	25
2.2.7.5.	Validation Level	25
2.2.7.5.1.	Domain Validation (DV)	25
2.2.7.5.2.	Organization Validation (OV)	25
2.2.7.5.3.	Extended Validation (EV)	26

2.2.7.6.	Identifier Validation Challenges	26
2.2.7.6.1.	HTTP Challenge	26
2.2.7.6.2.	DNS Challenge	27
2.3.	Protocolos asociados al Correo electrónico (explicacion de todo lo que hace un servidor de correo, y sus protocolos)	28
2.3.1.	¿Que es el protocolo SMTP?	28
2.3.2.	Recorrido completo de un mail	28
2.3.3.	SMTP con SSL	28
2.4.	Protocolos asociados a la consulta de un sitio (DNS)	28
2.4.1.	¿Que es el protocolo DNS?	28
2.4.2.	Recorrido completo de un mail	28
2.4.3.	SMTP con SSL	28
3.	(Nuevo) Herramientas a Utilizar	1
3.1.	El problema de los protocolos http y smtp	1
3.1.1.	Establishing Authority	1
3.1.2.	Risks of Intermediaries	2
3.1.3.	Attacks via Protocol Element Length	2
3.1.4.	Response Splitting	3
3.1.5.	Request Smuggling	3
3.1.6.	Message Integrity	4
3.1.7.	Message Confidentiality	4
3.2.	Network Security Attacks: basic concepts	4
3.2.1.	Passive Attacks	5
3.2.2.	Active Attacks	5
3.3.	Herramientas utilizadas	6
3.3.1.	Kali Linux	6
3.3.2.	Wireshark	7
3.3.3.	Ettrecap	8
3.4.	Virtualización	8
3.4.1.	Máquinas virtuales	8
3.4.2.	Container-Based Virtualization	9
3.4.3.	Docker	11

3.4.3.1.	Docker components	12
3.4.3.1.1.	Docker client and server	12
3.4.3.1.2.	Docker images	12
3.4.3.1.3.	Registries	12
3.4.3.1.4.	Containers	13
3.5.	Casos de estudio	14
3.5.1.	Caso de estudio: Sniffing de la red para obtener credenciales .	14
3.5.2.	Caso de estudio: Sniffing de la red para obtener mails inter- nos/externos	14
4.	Casos de estudio (Explotaciones y soluciones)	1
4.1.	Mejorando la seguridad en la navegación - Alternativas	1
4.1.1.	Self-signed Certificates	1
4.1.2.	Internal CA	1
4.1.3.	Estrategia utilizadas, ojala que con let's encrypt	1
4.2.	Estrategia	1
4.3.	CertBot para redes internas	1
4.4.	Mejorando la seguridad en los servidores de correo	1
5.	Conclusiones y Resultados Obtenidos	3
A.	Glosario	5
A.1.	Terminología	5
A.2.	Simbología	5

Capítulo 1

Introducción

1.1. Objetivos

Hence, not only the nodes but the communication channel should also be secured. While developing a secure network, Confidentiality, Integrity, Availability (CIA) needs to be considered

- item 1
- item 2

Este es un bien ambiente para
poner codigo

casa
casa
casa
casa
casa

Figura 1.1: Esta es la figura del escudo de la uns

1.2. Plan de tesis y principales contribuciones

- concientizar la implementacion de medidas de seguridad en redes internas
- demostrar como es posible realizar realizar acciones que perjudiquen a una organizacion
- implementar una mejora en la navegacion en una red interna utilizando docker

no se si las contribuciones se refiere a la principal fuente de consulta

1.3. Trabajos previos relacionados

Capítulo 2

(Ajustar) ¿Qué circula por una red interna?

2.1. Introducción

Network Technology is the key technology for a wide variety of applications such as email, file transfer, web browsing, online transactions, form fill up for various governmental or private activities, cab booking, etc. However, there is a significant lack of easy implementation of security methods for these applications.

2.2. Protocolos asociados a la web

2.2.1. ¿Que es el protocolo HTTP?

The Hypertext Transfer Protocol (HTTP) is a stateless application-level request/response protocol that uses extensible semantics and self-descriptive message payloads for flexible interaction with network-based hypertext information systems. HTTP is a generic interface protocol for information systems. It is designed to hide the details of how a service is implemented by presenting a uniform interface to clients that is independent of the types of resources provided. Likewise, servers do not need to be aware of each client's purpose: an HTTP request can be considered in isolation rather than being associated with a specific type of client or a predetermined sequence of application steps. The result is a protocol that can be

used effectively in many different contexts and for which implementations can evolve independently over time.

One consequence of this flexibility is that the protocol cannot be defined in terms of what occurs behind the interface. Instead, we are limited to defining the syntax of communication, the intent of received communication, and the expected behavior of recipients. If the communication is considered in isolation, then successful actions ought to be reflected in corresponding changes to the observable interface provided by servers. However, since multiple clients might act in parallel and perhaps at cross-purposes, we cannot require that such changes be observable beyond the scope of a single response.

2.2.2. Arquitectura

HTTP was created for the World Wide Web (WWW) architecture and has evolved over time to support the scalability needs of a worldwide hypertext system. Much of that architecture is reflected in the terminology and syntax productions used to define HTTP.

2.2.2.1. Client/Server Messaging

HTTP is a stateless request/response protocol that operates by exchanging messages across a reliable transport- or session-layer connection". An HTTP client is a program that establishes a connection to a server for the purpose of sending one or more HTTP requests. An HTTP server is a program that accepts connections in order to service HTTP requests by sending HTTP responses. Most HTTP communication consists of a retrieval request (GET) for a representation of some resource identified by a URI. In the simplest case, this might be accomplished via a single bidirectional connection (===) between the user agent (UA) and the origin server (O).

```
request -> UserAgent =====
Origin server |response
```

A client sends an HTTP request to a server in the form of a request message, beginning with a request-line that includes a method, URI, and protocol version, followed by header fields containing request modifiers, client information, and re-

presentation metadata, an empty line to indicate the end of the header section, and finally a message body containing the payload body (if any, Section 3.3). A server responds to a client's request by sending one or more HTTP response messages, each beginning with a status line that includes the protocol version, a success or error code, and textual reason phrase (Section 3.1.2), possibly followed by header fields containing server information, resource metadata, and representation metadata (Section 3.2), an empty line to indicate the end of the header section, and finally a message body containing the payload body (if any, Section 3.3).

2.2.2.2. Ejemplo

The following example illustrates a typical message exchange for a GET request (Section 4.3.1 of [RFC7231]) on the URI "http://www.example.com/hello.txt":

Client request:

```
GET /hello.txt HTTP/1.1 User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l  
zlib/1.2.3 Host: www.example.com Accept-Language: en, mi
```

Server response:

```
HTTP/1.1 200 OK Date: Mon, 27 Jul 2009 12:28:53 GMT Server: Apache  
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT ETag: "34aa387-d-1568eb00" Accept-  
Ranges: bytes Content-Length: 51 Vary: Accept-Encoding Content-Type: text/plain  
Hello World! My payload includes a trailing CRLF.
```

HTTP is defined as a stateless protocol, meaning that each request message can be understood in isolation. Many implementations depend on HTTP's stateless design in order to reuse proxied connections or dynamically load balance requests across multiple servers. Hence, a server **MUST NOT** assume that two requests on the same connection are from the same user agent unless the connection is secured and specific to that agent. Some non-standard HTTP extensions (e.g., [RFC4559]) have been known to violate this requirement, resulting in security and interoperability problems.

2.2.3. Formato del mensaje (Mejorar intro)

All HTTP/1.1 messages consist of a start-line followed by a sequence of octets in a format similar to the Internet Message Format [RFC5322]: zero or more header

fields (collectively referred to as the "headers" or the "header section"), an empty line indicating the end of the header section, and an optional message body. The normal procedure for parsing an HTTP message is to read the start-line into a structure, read each header field into a hash table by field name until the empty line, and then use the parsed data to determine if a message body is expected. If a message body has been indicated, then it is read as a stream until an amount of octets equal to the message body length is read or the connection is closed.

2.2.3.1. Start Line

An HTTP message can be either a request from client to server or a response from server to client. Syntactically, the two types of message differ only in the start-line, which is either a request-line (for requests) or a status-line (for responses), and in the algorithm for determining the length of the message body. In theory, a client could receive requests and a server could receive responses, distinguishing them by their different start-line formats, but, in practice, servers are implemented to only expect a request (a response is interpreted as an unknown or invalid request method) and clients are implemented to only expect a response.

2.2.3.1.1. Request Line A request-line begins with a method token, followed by a single space (SP), the request-target, another single space (SP), the protocol version, and ends with CRLF.

(podría ir un gráfico)

The method token indicates the request method to be performed on the target resource. The request method is case-sensitive.

The request-target identifies the target resource upon which to apply the request

2.2.3.1.2. Status Line The first line of a response message is the status-line, consisting of the protocol version, a space (SP), the status code, another space, a possibly empty textual phrase describing the status code, and ending with CRLF.

(podría ir un gráfico)

The status-code element is a 3-digit integer code describing the result of the server's attempt to understand and satisfy the client's corresponding request. The

rest of the response message is to be interpreted in light of the semantics defined for that status code.

2.2.3.2. Header Fields

Each header field consists of a case-insensitive field name followed by a colon (":"), optional leading whitespace, the field value, and optional trailing whitespace.

(podria ir un grafico o con un formato mejor) header-field = field-name ":" OWS field-value OWS

The field-name token labels the corresponding field-value as having the semantics defined by that header field.

2.2.3.2.1. Field Extensibility Header fields are fully extensible: there is no limit on the introduction of new field names, each presumably defining new semantics, nor on the number of header fields used in a given message. Existing fields are defined in each part of this specification and in many other specifications outside this document set.

New header fields can be defined such that, when they are understood by a recipient, they might override or enhance the interpretation of previously defined header fields, define preconditions on request evaluation, or refine the meaning of responses.

2.2.3.2.2. Field Order The order in which header fields with differing field names are received is not significant. However, it is good practice to send header fields that contain control data first, such as Host on requests and Date on responses, so that implementations can decide when not to handle a message as early as possible.

2.2.3.2.3. Whitespace This specification uses three rules to denote the use of linear whitespace: OWS (optional whitespace), RWS (required whitespace), and BWS ("bad" whitespace).

The OWS rule is used where zero or more linear whitespace octets might appear. The RWS rule is used when at least one linear whitespace octet is required to separate field tokens.

The BWS rule is used where the grammar allows optional whitespace only for historical reasons.

2.2.3.2.4. Field Parsing Messages are parsed using a generic algorithm, independent of the individual header field names. The contents within a given field value are not parsed until a later stage of message interpretation (usually after the message's entire header section has been processed). Consequently, this specification does not use ABNF rules to define each "Field-Name: Field Value" pair, as was done in previous editions. Instead, this specification uses ABNF rules that are named according to each registered field name, wherein the rule defines the valid grammar for that field's corresponding field values (i.e., after the field-value has been extracted from the header section by a generic field parser).

No whitespace is allowed between the header field-name and colon. In the past, differences in the handling of such whitespace have led to security vulnerabilities in request routing and response handling.

2.2.3.2.5. Field Limits HTTP does not place a predefined limit on the length of each header field or on the length of the header section as a whole. Various ad hoc limitations on individual header field length are found in practice, often depending on the specific field semantics.

2.2.3.2.6. Field Value Components (no entiendo, tal vez se puede sacar) Most HTTP header field values are defined using common syntax components (token, quoted-string, and comment) separated by whitespace or specific delimiting characters. Delimiters are chosen from the set of US-ASCII visual characters not allowed in a token

token = 1*tchar

/ DIGIT / ALPHA ; any VCHAR, except delimiters

A string of text is parsed as a single value if it is quoted using double-quote marks.

quoted-string = DQUOTE *(qdtext / quoted-pair) DQUOTE qdtext = HTAB / SP / obs-text =

Comments can be included in some HTTP header fields by surrounding the

comment text with parentheses. Comments are only allowed in fields containing comment.^{as} part of their field value definition.

comment = "("*(ctext / quoted-pair / comment) ")"
 ctext = HTAB / SP /

The backslash octet ("\") can be used as a single-octet quoting mechanism within quoted-string and comment constructs. Recipients that process the value of a quoted-string MUST handle a quoted-pair as if it were replaced by the octet following the backslash.

quoted-pair = "(" HTAB / SP / VCHAR / obs-text)

A sender SHOULD NOT generate a quoted-pair in a quoted-string except where necessary to quote DQUOTE and backslash octets occurring within that string. A sender SHOULD NOT generate a quoted-pair in a comment except where necessary to quote parentheses ["(.and ")"] and backslash octets occurring within that comment.

2.2.3.3. Message Body

The message body (if any) of an HTTP message is used to carry the payload body of that request or response. The message body is identical to the payload body unless a transfer coding has been applied.

The rules for when a message body is allowed in a message differ for requests and responses.

The presence of a message body in a request is signaled by a Content-Length or Transfer-Encoding header field. Request message framing is independent of method semantics, even if the method does not define any use for a message body.

The presence of a message body in a response depends on both the request method to which it is responding and the response status code. Responses to the HEAD request method never include a message body because the associated response header fields, if present, indicate only what their values would have been if the request method had been GET (Section 4.3.1 of [RFC7231]). 2xx (Successful) responses to a CONNECT request method (Section 4.3.6 of [RFC7231]) switch to tunnel mode instead of having a message body. All 1xx (Informational), 204 (No Content), and 304 (Not Modified) responses do not include a message body. All other responses do include a message body, although the body might be of zero

length.

2.2.3.3.1. Transfer-Encoding The Transfer-Encoding header field lists the transfer coding names corresponding to the sequence of transfer codings that have been (or will be) applied to the payload body in order to form the message body. Transfer codings are defined in Section 4.

Transfer-Encoding is analogous to the Content-Transfer-Encoding field of MIME, which was designed to enable safe transport of binary data over a 7-bit transport service . However, safe transport has a different focus for an 8bit-clean transfer protocol. In HTTP's case, Transfer-Encoding is primarily intended to accurately delimit a dynamically generated payload and to distinguish payload encodings that are only applied for transport efficiency or security from those that are characteristics of the selected resource. .

For example,

Transfer-Encoding: gzip, chunked

indicates that the payload body has been compressed using the gzip coding and then chunked using the chunked coding while forming the message body.

Unlike Content-Encoding , Transfer-Encoding is a property of the message, not of the representation, and any recipient along the request/response chain MAY decode the received transfer coding(s) or apply additional transfer coding(s) to the message body, assuming that corresponding changes are made to the Transfer-Encoding field-value. Additional information about the encoding parameters can be provided by other header fields not defined by this specification.

Transfer-Encoding was added in HTTP/1.1. It is generally assumed that implementations advertising only HTTP/1.0 support will not understand how to process a transfer-encoded payload.

2.2.3.3.2. Content-Length When a message does not have a Transfer-Encoding header field, a Content-Length header field can provide the anticipated size, as a decimal number of octets, for a potential payload body. For messages that do include a payload body, the Content-Length field-value provides the framing information necessary for determining where the body (and message) ends. For messages that do not include a payload body, the Content-Length indicates the size of the selected

representation .

2.2.3.3.3. Message Body Length (no entiendo, ver diferencia con content-length) The length of a message body is determined by one of the following (in order of precedence):

1. Any response to a HEAD request and any response with a 1xx (Informational), 204 (No Content), or 304 (Not Modified) status code is always terminated by the first empty line after the header fields, regardless of the header fields present in the message, and thus cannot contain a message body.

2. Any 2xx (Successful) response to a CONNECT request implies that the connection will become a tunnel immediately after the empty line that concludes the header fields. A client MUST ignore any Content-Length or Transfer-Encoding header fields received in such a message.

3. If a Transfer-Encoding header field is present and the chunked transfer coding (Section 4.1) is the final encoding, the message body length is determined by reading and decoding the chunked data until the transfer coding indicates the data is complete.

If a Transfer-Encoding header field is present in a response and the chunked transfer coding is not the final encoding, the message body length is determined by reading the connection until it is closed by the server. If a Transfer-Encoding header field is present in a request and the chunked transfer coding is not the final encoding, the message body length cannot be determined reliably; the server MUST respond with the 400 (Bad Request) status code and then close the connection.

If a message is received with both a Transfer-Encoding and a Content-Length header field, the Transfer-Encoding overrides the Content-Length. Such a message might indicate an attempt to perform request smuggling (Section 9.5) or response splitting (Section 9.4) and ought to be handled as an error. A sender MUST remove the received Content-Length field prior to forwarding such a message downstream.

4. If a message is received without Transfer-Encoding and with either multiple Content-Length header fields having differing field-values or a single Content-Length header field having an invalid value, then the message framing is invalid and the recipient MUST treat it as an unrecoverable error. If this is a request message, the server MUST respond with a 400 (Bad Request) status code and then close

the connection. If this is a response message received by a proxy, the proxy **MUST** close the connection to the server, discard the received response, and send a 502 (Bad Gateway) response to the client. If this is a response message received by a user agent, the user agent **MUST** close the connection to the server and discard the received response.

5. If a valid Content-Length header field is present without Transfer-Encoding, its decimal value defines the expected message body length in octets. If the sender closes the connection or the recipient times out before the indicated number of octets are received, the recipient **MUST** consider the message to be incomplete and close the connection.

6. If this is a request message and none of the above are true, then the message body length is zero (no message body is present).

7. Otherwise, this is a response message without a declared message body length, so the message body length is determined by the number of octets received prior to the server closing the connection.

Since there is no way to distinguish a successfully completed, close-delimited message from a partially received message interrupted by network failure, a server **SHOULD** generate encoding or length-delimited messages whenever possible. The close-delimiting feature exists primarily for backwards compatibility with HTTP/1.0.

A server **MAY** reject a request that contains a message body but not a Content-Length by responding with 411 (Length Required).

Unless a transfer coding other than chunked has been applied, a client that sends a request containing a message body **SHOULD** use a valid Content-Length header field if the message body length is known in advance, rather than the chunked transfer coding, since some existing services respond to chunked with a 411 (Length Required) status code even though they understand the chunked transfer coding. This is typically because such services are implemented via a gateway that requires a content-length in advance of being called and the server is unable or unwilling to buffer the entire request before processing.

A user agent that sends a request containing a message body **MUST** send a valid Content-Length header field if it does not know the server will handle HTTP/1.1 (or later) requests; such knowledge can be in the form of specific user configuration or by remembering the version of a prior received response.

If the final response to the last request on a connection has been completely received and there remains additional data to read, a user agent MAY discard the remaining data or attempt to determine if that data belongs as part of the prior response body, which might be the case if the prior message's Content-Length value is incorrect. A client MUST NOT process, cache, or forward such extra data as a separate response, since such behavior would be vulnerable to cache poisoning.

2.2.4. Métodos del protocolo HTTP

4.3.1. GET

The GET method requests transfer of a current selected representation for the target resource. GET is the primary mechanism of information retrieval and the focus of almost all performance optimizations. Hence, when people speak of retrieving some identifiable information via HTTP, they are generally referring to making a GET request.

It is tempting to think of resource identifiers as remote file system pathnames and of representations as being a copy of the contents of such files. In fact, that is how many resources are implemented. However, there are no such limitations in practice. The HTTP interface for a resource is just as likely to be implemented as a tree of content objects, a programmatic view on various database records, or a gateway to other information systems. Even when the URI mapping mechanism is tied to a file system, an origin server might be configured to execute the files with the request as input and send the output as the representation rather than transfer the files directly. Regardless, only the origin server needs to know how each of its resource identifiers corresponds to an implementation and how each implementation manages to select and send a current representation of the target resource in a response to GET.

4.3.2. HEAD

The HEAD method is identical to GET except that the server MUST NOT send a message body in the response (i.e., the response terminates at the end of the header section). The server SHOULD send the same header fields in response to a HEAD request as it would have sent if the request had been a GET, except that the payload header fields MAY be omitted. This method can be used for obtaining

metadata about the selected representation without transferring the representation data and is often used for testing hypertext links for validity, accessibility, and recent modification.

4.3.3. POST

The POST method requests that the target resource process the representation enclosed in the request according to the resource's own specific semantics. For example, POST is used for the following functions (among others):

- o Providing a block of data, such as the fields entered into an HTML form, to a data-handling process;
- o Posting a message to a bulletin board, newsgroup, mailing list, blog, or similar group of articles;
- o Creating a new resource that has yet to be identified by the origin server; and
- o Appending data to a resource's existing representation(s).

4.3.4. PUT

The PUT method requests that the state of the target resource be created or replaced with the state defined by the representation enclosed in the request message payload. A successful PUT of a given representation would suggest that a subsequent GET on that same target resource will result in an equivalent representation being sent in a 200 (OK) response. However, there is no guarantee that such a state change will be observable, since the target resource might be acted upon by other user agents in parallel, or might be subject to dynamic processing by the origin server, before any subsequent GET is received. A successful response only implies that the user agent's intent was achieved at the time of its processing by the origin server.

4.3.5. DELETE

The DELETE method requests that the origin server remove the association between the target resource and its current functionality. In effect, this method is similar to the `rm` command in UNIX: it expresses a deletion operation on the URI mapping of the origin server rather than an expectation that the previously associated information be deleted.

Relatively few resources allow the DELETE method – its primary use is for remote authoring environments, where the user has some direction regarding its effect. For example, a resource that was previously created using a PUT request, or identified via the Location header field after a 201 (Created) response to a POST request, might allow a corresponding DELETE request to undo those actions. Si-

milarly, custom user agent implementations that implement an authoring function, such as revision control clients using HTTP for remote operations, might use DELETE based on an assumption that the server's URI space has been crafted to correspond to a version repository.

4.3.6. CONNECT

The CONNECT method requests that the recipient establish a tunnel to the destination origin server identified by the request-target and, if successful, thereafter restrict its behavior to blind forwarding of packets, in both directions, until the tunnel is closed. Tunnels are commonly used to create an end-to-end virtual connection, through one or more proxies, which can then be secured using TLS (Transport Layer Security,).

CONNECT is intended only for use in requests to a proxy. However, most origin servers do not implement CONNECT.

There are significant risks in establishing a tunnel to arbitrary servers, particularly when the destination is a well-known or reserved TCP port that is not intended for Web traffic. For example, a CONNECT to a request-target of ".example.com:25" would suggest that the proxy connect to the reserved port for SMTP traffic; if allowed, that could trick the proxy into relaying spam email. Proxies that support CONNECT SHOULD restrict its use to a limited set of known ports or a configurable whitelist of safe request targets.

4.3.7. OPTIONS

The OPTIONS method requests information about the communication options available for the target resource, at either the origin server or an intervening intermediary. This method allows a client to determine the options and/or requirements associated with a resource, or the capabilities of a server, without implying a resource action. An OPTIONS request with an asterisk ("*") as the request-target applies to the server in general rather than to a specific resource. Since a server's communication options typically depend on the resource, the "*"request is only useful as a "ping" or "no-op" type of method; it does nothing beyond allowing the client to test the capabilities of the server. For example, this can be used to test a proxy for HTTP/1.1 conformance (or lack thereof).

A server generating a successful response to OPTIONS SHOULD send any header fields that might indicate optional features implemented by the server and appli-

cable to the target resource (e.g., Allow), including potential extensions not defined by this specification. The response payload, if any, might also describe the communication options in a machine or human-readable representation. A standard format for such a representation is not defined by this specification, but might be defined by future extensions to HTTP. A server **MUST** generate a Content-Length field with a value of "0" if no payload body is to be sent in the response.

4.3.8. TRACE

The TRACE method requests a remote, application-level loop-back of the request message. The final recipient of the request **SHOULD** reflect the message received, excluding some fields described below, back to the client as the message body of a 200 (OK) response with a Content-Type of "message/http". The final recipient is either the origin server or the first server to receive a Max-Forwards value of zero (0) in the request .

TRACE allows the client to see what is being received at the other end of the request chain and use that data for testing or diagnostic information. The value of the Via header field is of particular interest, since it acts as a trace of the request chain. Use of the Max-Forwards header field allows the client to limit the length of the request chain, which is useful for testing a chain of proxies forwarding messages in an infinite loop.

2.2.5. Response Status Codes

The status-code element is a three-digit integer code giving the result of the attempt to understand and satisfy the request.

HTTP status codes are extensible. HTTP clients are not required to understand the meaning of all registered status codes, though such understanding is obviously desirable. However, a client **MUST** understand the class of any status code, as indicated by the first digit, and treat an unrecognized status code as being equivalent to the x00 status code of that class, with the exception that a recipient **MUST NOT** cache a response with an unrecognized status code.

For example, if an unrecognized status code of 471 is received by a client, the client can assume that there was something wrong with its request and treat the response as if it had received a 400 (Bad Request) status code. The response message

will usually contain a representation that explains the status.

The first digit of the status-code defines the class of response. The last two digits do not have any categorization role. There are five values for the first digit:

- o 1xx (Informational): The request was received, continuing process
- o 2xx (Successful): The request was successfully received, understood, and accepted
- o 3xx (Redirection): Further action needs to be taken in order to complete the request
- o 4xx (Client Error): The request contains bad syntax or cannot be fulfilled
- o 5xx (Server Error): The server failed to fulfill an apparently valid request

2.2.6. Overview of Status Codes

The status codes listed below are defined in this specification, Section 4 of [RFC7232], Section 4 of [RFC7233], and Section 3 of [RFC7235]. The reason phrases listed here are only recommendations – they can be replaced by local equivalents without affecting the protocol.

Responses with status codes that are defined as cacheable by default (e.g., 200, 203, 204, 206, 300, 301, 404, 405, 410, 414, and 501 in this specification) can be reused by a cache with heuristic expiration unless otherwise indicated by the method definition or explicit cache controls [RFC7234]; all other status codes are not cacheable by default.

Code	Reason-Phrase
100	Continue
101	Switching Protocols
200	OK
201	Created
202	Accepted
203	Non-Authoritative Information
204	No Content
205	Reset Content
206	Partial Content

300	Multiple Choices
301	Moved Permanently
302	Found
303	See Other
304	Not Modified
305	Use Proxy
307	Temporary Redirect
400	Bad Request
401	Unauthorized
402	Payment Required
403	Forbidden
404	Not Found
405	Method Not Allowed
406	Not Acceptable
407	Proxy Authentication Required
408	Request Timeout
409	Conflict
410	Gone
411	Length Required
412	Precondition Failed
413	Payload Too Large
414	URI Too Long
415	Unsupported Media Type
416	Range Not Satisfiable
417	Expectation Failed
426	Upgrade Required
500	Internal Server Error
501	Not Implemented
502	Bad Gateway
503	Service Unavailable
504	Gateway Timeout

505	HTTP Version Not Supported
-----	----------------------------

2.2.6.1. Informational 1xx

The 1xx (Informational) class of status code indicates an interim response for communicating connection status or request progress prior to completing the requested action and sending a final response. 1xx responses are terminated by the first empty line after the status-line (the empty line signaling the end of the header section). Since HTTP/1.0 did not define any 1xx status codes, a server **MUST NOT** send a 1xx response to an HTTP/1.0 client.

2.2.6.2. Successful 2xx

The 2xx (Successful) class of status code indicates that the client's request was successfully received, understood, and accepted.

2.2.6.3. Redirection 3xx

The 3xx (Redirection) class of status code indicates that further action needs to be taken by the user agent in order to fulfill the request. If a Location header field (Section 7.1.2) is provided, the user agent **MAY** automatically redirect its request to the URI referenced by the Location field value, even if the specific status code is not understood. Automatic redirection needs to be done with care for methods not known to be safe, as defined in Section 4.2.1, since the user might not wish to redirect an unsafe request.

There are several types of redirects:

1. Redirects that indicate the resource might be available at a different URI, as provided by the Location field, as in the status codes 301 (Moved Permanently), 302 (Found), and 307 (Temporary Redirect).
2. Redirection that offers a choice of matching resources, each capable of representing the original request target, as in the 300 (Multiple Choices) status code.
3. Redirection to a different resource, identified by the Location field, that can represent an indirect response to the request, as in the 303 (See Other) status code.

4. Redirection to a previously cached result, as in the 304 (Not Modified) status code.

2.2.6.4. Client Error 4xx

The 4xx (Client Error) class of status code indicates that the client seems to have erred. Except when responding to a HEAD request, the server **SHOULD** send a representation containing an explanation of the error situation, and whether it is a temporary or permanent condition. These status codes are applicable to any request method. User agents **SHOULD** display any included representation to the user.

2.2.6.5. Server Error 5xx

The 5xx (Server Error) class of status code indicates that the server is aware that it has erred or is incapable of performing the requested method. Except when responding to a HEAD request, the server **SHOULD** send a representation containing an explanation of the error situation, and whether it is a temporary or permanent condition. A user agent **SHOULD** display any included representation to the user. These response codes are applicable to any request method.

2.2.7. HTTPS con SSL

With an understanding of some of the key concepts of cryptography, we can now look closely at the operation of the Secure Sockets Layer (ssl) protocol. Although ssl is not an extremely complicated protocol, it does offer several options and variations. The ssl protocol consists of a set of messages and rules about when to send (and not to send) each one. In this chapter, we consider what those messages are, the general information they contain, and how systems use the different messages in a communications session.

2.2.7.1. SSL Roles

The Secure Sockets Layer protocol defines two different roles for the communicating parties. One system is always a client, while the other is a server. The distinction is very important, because ssl requires the two systems to behave very differently.

The client is the system that initiates the secure communications; the server responds to the client's request. In the most common use of ssl, secure Web browsing, the Web browser is the ssl client and the Web site is the ssl server. For ssl itself, the most important distinctions between clients and servers are their actions during the negotiation of security parameters. Since the client initiates a communication, it has the responsibility of proposing a set of ssl options to use for the exchange. The server selects from the client's proposed options, deciding what the two systems will actually use. Although the final decision rests with the server, the server can only choose from among those options that the client originally proposed.

2.2.7.2. SSL Messages

When ssl clients and servers communicate, they do so by exchanging ssl messages. this chapter will show how systems use these messages in their communications. The most basic function that an ssl client and server can perform is establishing a channel for encrypted communications. This section looks at these steps in more detail by considering each message in the exchange.

GRAFIQUITO DE LOS MENSAJES

2.2.7.2.1. ClientHello The ClientHello message starts the ssl communication between the two parties. The client uses this message to ask the server to begin negotiating security services by using ssl.

The Version field of the ClientHello message contains the highest version number of ssl that the client can support. The current ssl version is 3.0, and it is by far the most widely deployed on the Internet.

The RandomNumber field, as you might expect, contains a random number. This random value, along with a similar random value that the server creates, provides the seed for critical cryptographic calculations. The ssl specification suggests that four of this field's 32 bytes consist of the time and date.

The next field in the ClientHello message is SessionID. Although all ClientHello messages may include this field, in this example, the field is meaningless and would be empty. The CipherSuites field allows a client to list the various cryptographic services that the client can support, including exact algorithms and key sizes. The

server actually makes the final decision as to which cryptographic services will be used for the communication, but it is limited to choosing from this list.

The `CompressionMethods` field is, in theory, similar to the `CipherSuites` field. In it, the client may list all of the various data compression methods that it can support. Compression methods are an important part of ssl because encryption has significant consequences on the effectiveness of any data compression techniques. Encryption changes the mathematical properties of information in a way that makes data compression virtually impossible.

2.2.7.2.2. ServerHello When the server receives the `ClientHello` message, it responds with a `ServerHello`. The `Version` field is the first example of a server making a final decision for the communications. The `ClientHello`'s version simply identifies which ssl versions the client can support. The `ServerHello`'s version, on the other hand, determines the ssl version that the communication will use. The `RandomNumber` field of the `ServerHello` is essentially the same as in the `ClientHello`, though this random value is chosen by the server. Along with the client's value, this number seeds important cryptographic calculations. The `SessionID` field of a `ServerHello` may contain a value, unlike the `ClientHello`'s field just discussed. The value in this case uniquely identifies this particular ssl communication, or session. The main reason for explicitly identifying a particular ssl session is to refer to it again later. The `CipherSuite` field (note that the name is singular, not plural, as in the case of a `ClientHello`) determines the exact cryptographic parameters, specifically algorithms and key sizes, to be used for the session. The server must select a single cipher suite from among those listed by the client in its `ClientHello` message. The `CompressionMethod` field is also singular for a `ServerHello`. In theory, the server uses this field to identify the data compression to be used for the session. Again, the server must pick from among those listed in the `ClientHello`. Current ssl versions have not defined any compression methods, however, so this field has no practical utility.

2.2.7.2.3. ServerKeyExchange In this example, the server follows its `ServerHello` message with a `ServerKeyExchange` message. This message complements the `CipherSuite` field of the `ServerHello`. While the `CipherSuite` field indicates the

cryptographic algorithms and key sizes, this message contains the public key information itself. The exact format of the key information depends on the particular public key algorithm used. For the rsa algorithm, for example, the server includes the modulus and public exponent of the server's rsa public key. Note that the `ServerKeyExchange` message is transmitted without encryption, so that only public key information can be safely included within it. The client will use the server's public key to encrypt a session key, which the parties will use to actually encrypt the application data for the session.

2.2.7.2.4. ServerHelloDone The `ServerHelloDone` message tells the client that the server has finished with its initial negotiation messages. The message itself contains no other information, but it is important to the client, because once the client receives a `ServerHelloDone`, it can move to the next phase of establishing the secure communications.

2.2.7.2.5. ClientKeyExchange When the server has finished its part of the initial ssl negotiation, the client responds with a `ClientKeyExchange` message. Just as the `ServerKeyExchange` provides the key information for the server, the `ClientKeyExchange` tells the server the client's key information. In

this case, however, the key information is for the symmetric encryption algorithm both parties will use for the session. Furthermore, the information in the client's message is encrypted using the public key of the server. This encryption protects the key information as it traverses the network, and it allows the client to verify that the server truly possesses the private key corresponding to its public key. Otherwise, the server won't be able to decrypt this message. This operation is an important protection against an attacker that intercepts messages from a legitimate server and pretends to be that server by forwarding the messages to an unsuspecting client. Since a fake server won't know the real server's private key, it won't be able to decrypt the `ClientKeyExchange` message. Without the information in that message, communication between the two parties cannot succeed.

2.2.7.2.6. ChangeCipherSpec After the client sends key information in a `ClientKeyExchange` message, the preliminary ssl negotiation is complete. At that point,

the parties are ready to begin using the security services they have negotiated. The ssl protocol defines a special message— `ChangeCipherSpec`—to explicitly indicate that the security services should now be invoked. Since the transition to secured communication is critical, and both parties have to get it exactly right, the ssl specification is very precise in describing the process. First, it identifies the set of information that defines security services. That information includes a specific symmetric encryption algorithm, a specific message integrity algorithm, and specific key material for those algorithms. The ssl specification also recognizes that some of that information (in particular, the key material) will be different for each direction of communication. In other words, one set of keys will secure data the client sends to the server, and a different set of keys will secure data the server sends to the client. (In principle, the actual algorithms could differ as well, but ssl does not define a way to negotiate such an option.) For any given system, whether it is a client or a server, ssl defines a write state and a read state. The write state defines the security information

for data that the system sends, and the read state defines the security information for data that the system receives. The `ChangeCipherSpec` message serves as the cue for a system to begin using its security information. Before a client or server sends a `ChangeCipherSpec` message, it must know the complete security information it is about to activate. As soon as the system sends this message, it activates its write state. Similarly, as soon as a system receives a `ChangeCipherSpec` from its peer, the system activates its read state.

GRAFIQUITO Figures 3-2 and 3-3

2.2.7.2.7. Finished Immediately after sending their `ChangeCipherSpec` messages, each system also sends a `Finished` message. The `Finished` messages allow both systems to verify that the negotiation has been successful and that security has not been compromised. Two aspects of the `Finished` message contribute to this security. First, as the previous subsection explained, the `Finished` message itself is subject to the negotiated cipher suite. That means that it is encrypted and authenticated according to that suite. If the receiving party cannot successfully decrypt and verify the message, then clearly something has gone awry with the security negotiation. The contents of the `Finished` message also serve to protect the security of the ssl

negotiation. Each Finished message contains a cryptographic hash of important information about the just-finished negotiation. Notice that protected data includes the exact content of all handshake messages used in the exchange (though ChangeCipherSpec messages are not considered “handshake” messages in the strict sense of the word, and thus are not included). This protects against an attacker who manages to insert fictitious messages or remove legitimate messages from the communication. If an attacker were able to do so, the client’s and server’s hash calculations would not match, and they would detect the compromise.

2.2.7.2.8. Ending Secure Communications Although as a practical matter it is rarely used (primarily due to the nature of Web sessions), ssl does have a defined procedure for ending a secure communication between two parties. In this procedure, the two systems each send a special ClosureAlert to the other. Explicitly closing a session protects against a truncation attack, in which an attacker is able to compromise security by prematurely terminating a communication.

2.2.7.3. Authenticating the Server’s Identity

previously it was explained how ssl can establish encrypted communications between two parties, that may not really add that much security to the communication. With encryption alone neither party can really be sure of the other’s identity. The typical reason for using encryption in the first place is to keep information secret from some third party. But if that third party were able to successfully masquerade as the intended recipient of the information, then encryption would serve no purpose. The data would be encrypted, but the attacker would have all the data necessary to decrypt it. To avoid this type of attack, ssl includes mechanisms that allow each party to authenticate the identity of the other. With these mechanisms, each party can be sure that the other is genuine, and not a masquerading attacker. In this section, we’ll look at how ssl enables a server to authenticate itself. A natural question is, of course, if authenticating identities is so important, why don’t we always authenticate both parties? Aca un ejemplo que sirva en una red interna The answer lies in the nature of Web commerce. When you want to purchase something using your Web browser, it’s very important that the Web site you’re browsing is authentic. You wouldn’t want to send your credit card number to some imposter

posing as your favorite merchant. The merchant, on the other hand, has other means for authenticating your identity. Once it receives a credit card number, for example, it can validate that number. Since the server doesn't need ssl to authenticate your identity, the ssl protocol allows for server authentication only.

2.2.7.3.1. Certificate When authenticating your identity, the server sends a Certificate message in place of the ServerKeyExchange message previously described. The Certificate message simply contains a certificate chain that begins with the server's public key certificate and ends with the certificate authority's root certificate. The client has the responsibility to make sure it can trust the certificate it receives from the server. That responsibility includes verifying the certificate signatures, validity times, and revocation status. It also means ensuring that the certificate authority is one that the client trusts. Typically, clients make this determination by knowing the public key of trusted certificate authorities in advance, through some trusted means. Netscape and Microsoft, for example, preload their browser software with public keys for well-known certificate authorities. Web servers that want to rely on this trust mechanism can only obtain their certificates (at least indirectly) from one of these wellknown authorities.

2.2.7.3.2. ClientKeyExchange The client's ClientKeyExchange message also differs in server authentication, though the difference is not major. When encryption only is to be used, the client encrypts the information in the ClientKeyExchange using the public key the server provides in its ServerKeyExchange message. In this case, of course, the server is authenticating itself and, thus, has sent a Certificate message instead of a ServerKeyExchange. The client, therefore, encrypts its ClientKeyExchange information using the public key contained in the server's certificate. This step is important because it allows the client to make sure that the party with whom it is communicating actually possesses the server's private key. Only a system with the actual private key will be able to decrypt this message and successfully continue the communication.

2.2.7.4. Certificate Functionality

2.2.7.4.1. Single Domain As the name suggests, a single domain SSL certificate can only be used on a single domain or IP. This is considered the default SSL certificate type. The DV SSL type is available at all validation levels.

2.2.7.4.2. Multi-Domain This SSL type is a jack-of-all-trades certificate. Multi-Domain Wildcards can encrypt up to 250 different domains and unlimited sub-domains. Unfortunately, it's not available in EV.

2.2.7.4.3. Wildcard Wildcards are specifically designed to encrypt one domain and all of its accompanying sub-domains. Unlimited sub-domains. Unfortunately, Wildcards are only available at the DV and OV levels.

2.2.7.4.4. Multi-Domain Wildcard These are the jack-of-all-trades certificates. Multi-Domain Wildcards can encrypt up to 250 different domains and unlimited sub-domains. Unfortunately, it's not available in EV.

2.2.7.5. Validation Level

There are three types of SSL Certificate available today; Extended Validation (EV SSL), Organization Validated (OV SSL) and Domain Validated (DV SSL). The encryption levels are the same for each certificate, what differs is the vetting and verification processes needed to obtain the certificate.

2.2.7.5.1. Domain Validation (DV) Domain Validation SSL or DV SSL represents the base-level for SSL types. These are perfect for websites that just need encryption and nothing more. DV SSL certificates are typically inexpensive and they can be issued within minutes. That's because the validation process is fully automated. Just prove you own your domain and the DV SSL certificate is yours.

2.2.7.5.2. Organization Validation (OV) Organization Validation SSL or OV SSL represents the middle ground for SSL certificate types. To obtain OV SSL, your company or organization must undergo light business vetting. This can

take up to three business days because someone has to verify your business information. OV SSL displays the same visual indicators as DV SSL, but provides a way for your customers to check your verified business information in the certificate details section.

2.2.7.5.3. Extended Validation (EV) Extended Validation SSL or EV SSL requires extensive business vetting by Comodo. This may sound like a lot, but it's really not if your business has publicly available records. EV SSL activates a unique visual indicator – your verified organization name shown in the browser.

2.2.7.6. Identifier Validation Challenges

(CAMBIAR LA INTRO, NO ME GUSTA LO DE IDENTIFICADOR, RELACIONARLO MAS CON UN DOMINIO)

ACME uses an extensible challenge/response framework for identifier validation. The server presents a set of challenges in the authorization object it sends to a client (as objects in the `challenges.array`), and the client responds by sending a response object in a POST request to a challenge URL.

Different challenges allow the server to obtain proof of different aspects of control over an identifier. In some challenges, like HTTP and DNS, the client directly proves its ability to do certain things related to the identifier. The choice of which challenges to offer to a client under which circumstances is a matter of server policy.

2.2.7.6.1. HTTP Challenge With HTTP validation, the client in an ACME transaction proves its control over a domain name by proving that it can provision HTTP resources on a server accessible under that domain name. The ACME server challenges the client to provision a file at a specific path, with a specific string as its content.

This is the most common challenge type today. The server gives a token to your ACME client, and your ACME client puts a file on your web server at `http://YOUR_DOMAIN/.well-known/acme-challenge/TOKEN`. That file contains the token, plus a thumbprint of your account key.

Once the client tells to the server that the file is ready, the server tries retrieving it. On receiving a response, the server constructs and stores the key authorization

from the challenge "token" value and the current client account key.

Given a challenge/response pair, the server verifies the client's control of the domain by verifying that the resource was provisioned as expected.

(TAL VEZ PARA LA PRESENTACION)

Pros:

It's easy to automate without extra knowledge about a domain's configuration. It allows hosting providers to issue certificates for domains CNAMEd to them. It works with off-the-shelf web servers.

Cons:

It doesn't work if your ISP blocks port 80 (this is rare, but some residential ISPs do this). Let's Encrypt doesn't let you use this challenge to issue wildcard certificates. If you have multiple web servers, you have to make sure the file is available on all of them.

(EXPLICAR POR QUE NO PUEDO USAR ESTE DESAFIO)

2.2.7.6.2. DNS Challenge When the identifier being validated is a domain name, the client can prove control of that domain by provisioning a TXT resource record containing a designated value for a specific validation domain name.

A client fulfills this challenge by constructing a key authorization from the "token" value provided in the challenge and the client's account key. The client then computes the SHA-256 digest of the key authorization.

The record provisioned to the DNS contains the base64url encoding of this digest. The client constructs the validation domain name by prepending the label `_acme-challenge` to the domain name being validated, then provisions a TXT record with the digest value under that name. For example, if the domain name being validated is "www.example.org", then the client would provision the following DNS record: `_acme-challenge.www.example.org. 300 IN TXT "gfj9Xq...Rg85nM"`

On receiving a response, the server constructs and stores the key authorization from the challenge "token" value and the current client account key.

To validate a DNS challenge, the server performs the following steps:

1. Compute the SHA-256 digest of the stored key authorization
2. Query for TXT records for the validation domain name
3. Verify that the contents of one of the TXT records match the digest value

If all of the above verifications succeed, then the validation is successful. If no DNS record is found, or DNS record and response payload do not pass these checks, then the validation fails.

The client SHOULD de-provision the resource record(s) provisioned for this challenge once the challenge is complete, i.e., once the "status" field of the challenge has the value "valid" or "invalid".

2.3. Protocolos asociados al Correo electrónico (explicación de todo lo que hace un servidor de correo, y sus protocolos)

2.3.1. ¿Que es el protocolo SMTP?

2.3.2. Recorrido completo de un mail

2.3.3. SMTP con SSL

2.4. Protocolos asociados a la consulta de un sitio (DNS)

2.4.1. ¿Que es el protocolo DNS?

2.4.2. Recorrido completo de un mail

2.4.3. SMTP con SSL

Capítulo 3

(Nuevo) Herramientas a Utilizar

En éste capítulo se verá por qué los protocolos http y smtp son inseguros, introducción al snnifin, spoofing, arp attack

3.1. El problema de los protocolos http y smtp

(Agregar una intro)

3.1.1. Establishing Authority

HTTP relies on the notion of an authoritative response: a response that has been determined by (or at the direction of) the authority identified within the target URI to be the most appropriate response for that request given the state of the target resource at the time of response message origination. Providing a response from a non-authoritative source, such as a shared cache, is often useful to improve performance and availability, but only to the extent that the source can be trusted or the distrusted response can be safely used.

Unfortunately, establishing authority can be difficult. For example, phishing is an attack on the user's perception of authority, where that perception can be misled by presenting similar branding in hypertext, possibly aided by userinfo obfuscating the authority component (see Section 2.7.1). User agents can reduce the impact of phishing attacks by enabling users to easily inspect a target URI prior to making an action, by prominently distinguishing (or rejecting) userinfo when present, and

by not sending stored credentials and cookies when the referring document is from an unknown or untrusted source.

When a registered name is used in the authority component, the "httpURI scheme (Section 2.7.1) relies on the user's local name resolution service to determine where it can find authoritative responses. This means that any attack on a user's network host table, cached names, or name resolution libraries becomes an avenue for attack on establishing authority. Likewise, the user's choice of server for Domain Name Service (DNS), and the hierarchy of servers from which it obtains resolution results, could impact the authenticity of address mappings; DNS Security Extensions are one way to improve authenticity.

Furthermore, after an IP address is obtained, establishing authority for an "httpURI is vulnerable to attacks on Internet Protocol routing.

3.1.2. Risks of Intermediaries

By their very nature, HTTP intermediaries are men-in-the-middle and, thus, represent an opportunity for man-in-the-middle attacks. Compromise of the systems on which the intermediaries run can result in serious security and privacy problems. Intermediaries might have access to security-related information, personal information about individual users and organizations, and proprietary information belonging to users and content providers. A compromised intermediary, or an intermediary implemented or configured without regard to security and privacy considerations, might be used in the commission of a wide range of potential attacks.

Intermediaries that contain a shared cache are especially vulnerable to cache poisoning attacks. Implementers need to consider the privacy and security implications of their design and coding decisions, and of the configuration options they provide to operators (especially the default configuration).

Users need to be aware that intermediaries are no more trustworthy than the people who run them; HTTP itself cannot solve this problem.

3.1.3. Attacks via Protocol Element Length

Because HTTP uses mostly textual, character-delimited fields, parsers are often vulnerable to attacks based on sending very long (or very slow) streams of data,

particularly where an implementation is expecting a protocol element with no pre-defined length.

3.1.4. Response Splitting

Response splitting (a.k.a, CRLF injection) is a common technique, used in various attacks on Web usage, that exploits the line-based nature of HTTP message framing and the ordered association of requests to responses on persistent connections [Klein]. This technique can be particularly damaging when the requests pass through a shared cache.

Response splitting exploits a vulnerability in servers (usually within an application server) where an attacker can send encoded data within some parameter of the request that is later decoded and echoed within any of the response header fields of the response. If the decoded data is crafted to look like the response has ended and a subsequent response has begun, the response has been split and the content within the apparent second response is controlled by the attacker. The attacker can then make any other request on the same persistent connection and trick the recipients (including intermediaries) into believing that the second half of the split is an authoritative answer to the second request.

For example, a parameter within the request-target might be read by an application server and reused within a redirect, resulting in the same parameter being echoed in the Location header field of the response. If the parameter is decoded by the application and not properly encoded when placed in the response field, the attacker can send encoded CRLF octets and other content that will make the application's single response look like two or more responses.

3.1.5. Request Smuggling

Request smuggling ([Linhart]) is a technique that exploits differences in protocol parsing among various recipients to hide additional requests (which might otherwise be blocked or disabled by policy) within an apparently harmless request. Like response splitting, request smuggling can lead to a variety of attacks on HTTP usage.

3.1.6. Message Integrity

HTTP does not define a specific mechanism for ensuring message integrity, instead relying on the error-detection ability of underlying transport protocols and the use of length or chunk-delimited framing to detect completeness. Additional integrity mechanisms, such as hash functions or digital signatures applied to the content, can be selectively added to messages via extensible metadata header fields. Historically, the lack of a single integrity mechanism has been justified by the informal nature of most HTTP communication. However, the prevalence of HTTP as an information access mechanism has resulted in its increasing use within environments where verification of message integrity is crucial.

User agents are encouraged to implement configurable means for detecting and reporting failures of message integrity such that those means can be enabled within environments for which integrity is necessary. For example, a browser being used to view medical history or drug interaction information needs to indicate to the user when such information is detected by the protocol to be incomplete, expired, or corrupted during transfer. Such mechanisms might be selectively enabled via user agent extensions or the presence of message integrity metadata in a response. At a minimum, user agents ought to provide some indication that allows a user to distinguish between a complete and incomplete response message (Section 3.4) when such verification is desired.

3.1.7. Message Confidentiality

HTTP relies on underlying transport protocols to provide message confidentiality when that is desired. HTTP has been specifically designed to be independent of the transport protocol, such that it can be used over many different forms of encrypted connection, with the selection of such transports being identified by the choice of URI scheme or within user agent configuration.

3.2. Network Security Attacks: basic concepts

There are mainly two types of network attacks – passive attack and active attack

Passive: This type of attack happens when sensitive information is monitored and analyzed, possibly compromising the security of enterprises and their customers. In short, network intruder intercepts data traveling through the network. • Active: This type of attack happens when information is modified, altered or demolished entirely by a hacker. Here the interloper starts instructions to disturb the network's regular process.

So the motives behind passive attackers and active attackers are totally different. Whereas the motive of passive attackers is simply to steal sensitive information and to analyze the traffic to steal future messages, the motive of active attackers is to stop normal communication between two legitimate entities.

3.2.1. Passive Attacks

Passive attackers are mainly interested in stealing sensitive information. This happens without the knowledge of the victim. As such passive attacks are difficult to detect and thereby secure the network. The following are some of the passive attacks that are in existence [7].

- Traffic Analysis: Attacker senses the communication path between the sender and receiver.
- Monitoring: Attacker can read the confidential data, but he cannot edit or modify the data.
- Eavesdropping: This type of attack occurs in the mobile ad-hoc network where basically the attacker finds out some secret or confidential information from communication.

3.2.2. Active Attacks

The active attacks happen in such a manner so as to notify the victims that their systems have been compromised. As a result, the victim stops communication with the other party. Some of the active attacks are as follows [8].

- Modification: Some alterations in the routing route is performed by the malicious node. This results in causing the sender to send messages through the

long route, which causes communication delay. This is an attack on integrity as shown in Fig. 2.

- Wormhole: This attack is also called a tunneling attack. A packet is received by an attacker at one point. He then tunnels it to another malicious node in the network. This causes a beginner to assume that he found the shortest path in the network as shown in Fig. 3.
- Fabrication: A malicious node generates a false routing message that causes the generation of incorrect information about the route between devices. This is an attack on authenticity as shown in Fig. 4.
- Spoofing: A malicious node miss-present his identity so that the sender changes his topology as shown in Fig. 5.
- Denial of services: A malicious node sends a message to the node and consumes the bandwidth of the network as given in Fig.
- Man-in-the-middle: Attack—Also called a hijacking attack, it is an attack where the attacker secretly alters and relays the communications between two legitimate parties without their knowledge. These parties in turn are unaware of the secret hacker consider that they are doing direct communication with each other [13]. Figure 12 depicts this attack.

3.3. Herramientas utilizadas

3.3.1. Kali Linux

BackTrack is one of the most famous Linux distribution systems, as can be proven by the number of downloads that reached more than four million as of BackTrack Linux 4.0 pre final.

Kali Linux Version 1.0 was released on March 12, 2013. Five days later, Version 1.0.1 was released, which fixed the USB keyboard issue. In those five days, Kali has been downloaded more than 90,000 times.

Kali Linux is security-focused Linux distribution based on Debian. It's a rebranded version of the famous Linux distribution known as Backtrack, which came with

a huge repository of open source hacking tools for network, wireless, and web application penetration testing. Although Kali Linux contains most of the tools from Backtrack, the main aim of Kali Linux is to make it portable so that it could be installed on devices based on the ARM architectures such as tablets and Chromebook, which makes the tools available at your disposal with much ease.

Using open source hacking tools comes with a major drawback: they contain a whole lot of dependencies when installed on Linux and they need to be installed in a predefined sequence. Moreover, authors of some tools have not released accurate documentation, which makes our life difficult.

Kali Linux simplifies this process; it contains many tools preinstalled with all the dependencies and is in ready to use condition so that you can pay more attention for the actual attack and not on installing the tool. Updates for tools installed in Kali Linux are more frequently released, which helps you to keep the tools up to date. A non-commercial toolkit that has all the major hacking tools preinstalled to test realworld networks and applications is a dream of every ethical hacker and the authors of Kali Linux make every effort to make our life easy, which enables us to spend more time on finding the actual flaws rather than building a toolkit.

(EXPLICAR COMO Y DONDE LO INSTALE) (DECIR QUE FUE EL SISTEMA OPERATIVO UTILIZADO QUE CONTIENEN LAS SIGUIENTES HERRAMIENTAS)

3.3.2. Wireshark

Wireshark is one of the most popular, free, and open source network protocol analyzers. Wireshark is preinstalled in Kali and ideal for network troubleshooting, analysis, and for this chapter, a perfect tool to monitor traffic from potential targets with the goal of capturing session tokens. Wireshark uses a GTK+ widget toolkit to implement its user interface and pcap to capture packets. It operates very similarly to a tcpdump command; however, acting as a graphical frontend with integrated sorting and filtering options. (HAY MAS, CON GRAFICOS EN) Joseph Muniz, Aamir Lakhani - Web Penetration Testing with Kali Linux-Packt Publishing (2013)

3.3.3. Ettrcap

Ettercap is a free and open source comprehensive suite for man-in-the-middle-based attacks. Ettercap can be used for computer network protocol analysis and security auditing, featuring sniffing live connections, content filtering, and support for active and passive dissection of multiple protocols. Ettercap works by putting the attacker's network interface into promiscuous mode and ARP for poisoning the victim machines.

3.4. Virtualización

(AGREGAR GRAFICOS SI O SI) Virtualization provides abstraction on top of the actual resources we want to virtualize. The level at which this abstraction is applied changes the way that different virtualization techniques look. At a higher level, there are two major virtualization techniques based on the level of abstraction.

- Virtual machine (VM)-based
- Container-based

Apart from these two virtualizing techniques, there are other techniques, such as unikernels, which are lightweight single-purpose VMs. IBM is currently attempting to run unikernels as processes with projects like Nabla. In this book, we will mainly look at VM-based and container-based virtualizations only.

3.4.1. Máquinas virtuales

VM-Based Virtualization The VM-based approach virtualizes the complete OS. The abstraction it presents to the VM are virtual devices like virtual disks, virtual CPUs, and virtual NICs. In other words, we can state that this is virtualizing the complete ISA (instruction set architecture); as an example, the x86 ISA. With virtual machines, multiple OSes can share the same hardware resources, with virtualized representations of each of the resources available to the VM. For example, the OS on the virtual machine (also called the guest) can continue to do I/O operations on a disk (in this case, it's a virtual disk), thinking that it's the only OS running on the physical hardware (also called the host), although in actuality, it is shared by multiple virtual machines as well as by the host OS. VMs are very similar to other processes in the host OS. VMs execute in a hardware-isolated virtual address space

and at a lower privilege level than the host OS. The primary difference between a process and a VM is the ABI (Application Binary Interface) exposed by the host to the VM. In the case of a process, the exposed ABI has constructs like network sockets, FDs, and so on, whereas with a full-fledged OS virtualization, the ABI will have a virtual disk, a virtual CPU, virtual network cards, and so on.

Hypervisors A special piece of software is used to virtualize the OS, called the hypervisor. The hypervisor itself has two parts: • **Virtual Machine Monitor (VMM)**: Used for trapping and emulating the privileged instruction set (which only the kernel of the operating system can perform). The VMM has to satisfy three properties (Popek and Goldberg, 1973): • **Isolation** : Should isolate guests (VMs) from each other. • **Equivalency** : Should behave the same, with or without virtualization. This means we run the majority (almost all) of the instructions on the physical hardware without any translation, and so on. • **Performance** : Should perform as good as it does without any virtualization. This again means that the overhead of running a VM is minimal.

Device Model The device model of the hypervisor handles the I/O virtualization again by trapping and emulating and then delivering interrupts back to the specific virtual machine. The device model handles: • **Memory Virtualization** • **Shadow Page Table** • **CPU Virtualization** • **IO Virtualization**

3.4.2. Container-Based Virtualization

Container-Based Virtualization This form of virtualization doesn't abstract the hardware but uses techniques within the Linux kernel to isolate access paths for different resources. It carves out a logical boundary within the same operating system. As an example, we get a separate root file system, a separate process tree, a separate network subsystem, and so on.

The Linux kernel is made up of several components and functionalities; the ones related to containers are as follows: **Control groups (cgroups)** **Namespaces** **Security-Enhanced Linux (SELinux)** **Cgroups** The cgroup functionality allows for limiting and prioritizing resources, such as CPUs, RAM, the network, the filesystem, and so on. The main goal is to not exceed the resources—to avoid wasting resources that might be needed for other processes. **Namespaces** The namespace functiona-

lity allows for partitioning of kernel resources, such that one set of processes sees one set of resources, while another set of processes sees a different set of resources. The feature works by having the same namespace for these resources in the various sets of processes, but having those names refer to distinct resources. Examples of resource names that can exist in multiple spaces (so that the named resources will be partitioned) are process IDs, hostnames, user IDs, filenames, and some names associated with network access and inter-process communication. When a Linux system boots; that is, only one namespace is created. Processes and resources will join the same namespace, until a different namespace is created, resources assigned to it, and processes join it. SELinux SELinux is a module of the Linux kernel that provides a mechanism to enforce the security of the system, with specific policies. Basically, SELinux can limit programs from accessing files and network resources. The idea is to limit the privileges of programs and daemons to a minimum, so that it can limit the risk of system halt. The preceding functionalities have been around for many years. Namespaces were first released in 2002, and cgroups in 2005, by Google (cgroups were first named process containers, and then cgroups). For example, SunSolaris 5.10, released at the beginning of 2005, provided support for Solaris containers. Nowadays, Linux containers are the new buzzword, and some people think they are a new means of virtualization. Containerization uses resources directly, and does not need an emulator at all; the fewer resources, the more efficiency. Different applications can run on the same host: isolated at the kernel level and isolated by namespaces and cgroups. The kernel (that is, the OS) is shared by all containers, as shown in the following diagram: Containers When we talk about containers, we are indirectly referring to two main concepts—a container image and a running container image. A container image is the definition of the container, wherein all software stacks are installed as additional layers, as depicted by the following diagram: A container image is typically made up of multiple layers. The first layer is given by the base image, which provides the OS core functionalities, with all of the tools needed to get started. Teams often work by building their own layers on these base images. Users can also build on more advanced application images, which not only have an OS, but which also include language runtimes, debugging tools, and libraries, as shown in the following diagram: Base images are built from the same utilities and libraries that are included in an OS. A good base image provides a secure and

stable platform on which to build applications. Red Hat provides base images for Red Hat Enterprise Linux. These images can be used like a normal OS. Users can customize them for their applications as necessary, installing packages and enabling services to start up just like a normal Red Hat Enterprise Linux Server. Containers provide isolation by taking advantage of kernel technologies, like cgroups, kernel namespaces, and SELinux, which have been battle-tested and used for years at Google and the US Department of Defense, in order to provide application isolation. Since containers use a shared kernel and container host, they reduce the amount of resources required for the container itself, and are more lightweight when compared to VMs. Therefore, containers provide an unmatched agility that is not feasible with VMs; for example, it only takes a few seconds to start a new container. Furthermore, containers support a more flexible model when it comes to CPU utilization and memory resources, and allow for resource burst modes, so that applications can consume more resources when required, within the defined boundaries.

3.4.3. Docker

Docker Docker is an open-source engine that automates the deployment of applications into containers. It was written by the team at Docker, Inc (formerly dotCloud Inc, an early player in the Platform-as-a-Service (PAAS) market), and released by them under the Apache 2.0 license. Docker adds an application deployment engine on top of a virtualized container execution environment. It is designed to provide a lightweight and fast environment in which to run your code as well as an efficient workflow to get that code from your laptop to your test environment and then into production. Docker is incredibly simple. Indeed, you can get started with Docker on a minimal host running nothing but a compatible Linux kernel and a Docker binary.

With Docker, Developers care about their applications running inside containers, and Operations cares about managing the containers. Docker is designed to enhance consistency by ensuring the environment in which your developers write code matches the environments into which your applications are deployed. This reduces the risk of “worked in dev, now an ops problem.”

3.4.3.1. Docker components

Let's look at the core components that compose the Docker Community Edition: The Docker client and server, also called the Docker Engine. Docker Images Registries Docker Containers

3.4.3.1.1. Docker client and server Docker is a client-server application. The Docker client talks to the Docker server or daemon, which, in turn, does all the work. You'll also sometimes see the Docker daemon called the Docker Engine. Docker ships with a command line client binary, `docker`, as well as a full RESTful API to interact with the daemon: `dockerd`. You can run the Docker daemon and client on the same host or connect your local Docker client to a remote daemon running on another host. You can see Docker's architecture depicted here:

3.4.3.1.2. Docker images Images are the building blocks of the Docker world. You launch your containers from images. Images are the "build" part of Docker's life cycle. They have a layered format, using Union file systems, that are built step-by-step using a series of instructions. For example: Add a file. Run a command. Open a port. You can consider images to be the "source code" for your containers. They are highly portable and can be shared, stored, and updated. In the book, we'll learn how to use existing images as well as build our own images.

3.4.3.1.3. Registries Docker stores the images you build in registries. There are two types of registries: public and private. Docker, Inc., operates the public registry for images, called the Docker Hub. You can create an account on the Docker Hub and use it to share and store your own images. The Docker Hub also contains, at last count, over 10,000 images that other people have built and shared. Want a Docker image for an Nginx web server, the Asterisk open source PABX system, or a MySQL database? All of these are available, along with a whole lot more. You can also store images that you want to keep private on the Docker Hub. These images might include source code or other proprietary information you want to keep secure or only share with other members of your team or organization. You can also run your own private registry, and we'll show you how to do that in Chapter 4. This

allows you to store images behind your firewall, which may be a requirement for some organizations.

3.4.3.1.4. Containers Docker helps you build and deploy containers inside of which you can package your applications and services. As we’ve just learned, containers are launched from images and can contain one or more running processes. You can think about images as the building or packing aspect of Docker and the containers as the running or execution aspect of Docker. A Docker container is: An image format. A set of standard operations. An execution environment. Docker borrows the concept of the standard shipping container, used to transport goods globally, as a model for its containers. But instead of shipping goods, Docker containers ship software. Each container contains a software image – its ‘cargo’ – and, like its physical counterpart, allows a set of operations to be performed. For example, it can be created, started, stopped, restarted, and destroyed. Like a shipping container, Docker doesn’t care about the contents of the container when performing these actions; for example, whether a container is a web server, a database, or an application server. Each container is loaded the same as any other container. Docker also doesn’t care where you ship your container: you can build on your laptop, upload to a registry, then download to a physical or virtual server, test, deploy to a cluster of a dozen Amazon EC2 hosts, and run. Like a normal shipping container, it is interchangeable, stackable, portable, and as generic as possible. With Docker, we can quickly build an application server, a message bus, a utility appliance, a CI test bed for an application, or one of a thousand other possible applications, services, and tools. It can build local, self-contained test environments or replicate complex application stacks for production or development purposes. The possible use cases are endless.

3.5. Casos de estudio

3.5.1. Caso de estudio: Sniffing de la red para obtener credenciales

(HAY MAS, CON GRAFICOS EN) Joseph Muniz, Aamir Lakhani - Web Penetration Testing with Kali Linux-Packt Publishing (2013) Aca yo segui un tutorial, buscarlo

3.5.2. Caso de estudio: Sniffing de la red para obtener mails internos/externos

Capítulo 4

Casos de estudio (Explotaciones y soluciones)

4.1. Mejorando la seguridad en la navegación - Alternativas

4.1.1. Self-signed Certificates

4.1.2. Internal CA

4.1.3. Estrategia utilizadas, ojala que con let's encrypt

4.2. Estrategia

4.3. CertBot para redes internas

4.4. Mejorando la seguridad en los servidores de correo

Capítulo 5

Conclusiones y Resultados Obtenidos

Apéndice A

Glosario

A.1. Terminología

Término en inglés	Traducción utilizada
argument	argumento
argumentative system	sistema argumentativo
assumption	suposición
atom	átomo
backing	fundamentos
blocking defeater	derrotador de bloqueo
burden of proof	peso de la prueba
claim	afirmación

A.2. Simbología

Símbolo	Página	Significado
$\neg h$	103	negación fuerte del átomo h

Bibliografía

- [1] BONDARENKO, A., DUNG, P. M., KOWALSKI, R., AND TONI, F. An abstract argumentation-theoretic approach to default reasoning. *Artificial Intelligence* 93, 1–2 (1997), 63–101.
- [2] CAPOBIANCO, M. El Rol de las Bases de Dialéctica en la Argumentación Rebatible. tesis de licenciatura, July 1999.
- [3] CAPOBIANCO, M., CHESÑEVAR, C. I., AND SIMARI, G. R. An argumentative formalism for implementing rational agents. In *Proceedings del 2do Workshop en Agentes y Sistemas Inteligentes (WASI), 7mo Congreso Argentino de Ciencias de la Computación (CACIC)* (El Calafate, Santa Cruz, Oct. 2001), Universidad Nacional de la Patagonia Austral, pp. 1051–1062.
- [4] CHESÑEVAR, C. I. *Formalización de los Procesos de Argumentación Rebatible como Sistemas Deductivos Etiquetados*. PhD thesis, Departamento de Ciencias de la Computación, Universidad Nacional del Sur, Bahía Blanca, Argentina, Jan. 2001.
- [5] DAVIS, R. E. *Truth, Deduction, and Computation*. Computer Science Press, 1989.
- [6] GARCÍA, A. J. La Programación en Lógica Rebatible: su definición teórica y computacional. Master’s thesis, Departamento de Ciencias de la Computación, Universidad Nacional del Sur, Bahía Blanca, Argentina, July 1997.
- [7] HAENNI, R. Modeling uncertainty with propositional assumption-based systems. In *Applications of uncertainty formalisms*, A. Hunter and S. Parsons, Eds. Springer-Verlag, 1998, pp. 446–470.